THE ART OF AI MONKS

Avatars, Attention, and the Business Behind the Robes



ANICCA

The Art of AI Monks

Avatars, Attention, and the Business Behind the Robes

Anicca

First edition. Research snapshot: 21–22 October 2026.

A note to the reader

A monk appears on your phone. He seems unhurried. You are holding a device that contains seven unanswered messages, a shopping basket, and a small machine for comparing your life with other people's lives. He would like one minute of your attention.

There is something funny about this encounter. There is also something serious. The same image can carry a thoughtful reflection, an unsupported health claim, an original story, a sales pitch, or several of those things at once. The person may be synthetic. The person may be real and using a digital character. The available evidence may not tell us.

This book asks how the format works, what public records establish about its business, and how to make something original from the useful parts. It is a cultural field guide with a working manual inside it.

You can read it for the strange spectacle of a virtual monastery with a checkout page. You can also use it to plan, produce, and measure an honest character of your own.

Our starting archive contains **75 records and 48 source references: nine researched profiles and 66 discovery leads**. Its AI-status fields contain four `reported`, 71 `unverified`, and zero `confirmed` records. "Researched" describes work performed on a profile; it does not certify its AI origin. This is an English-first sample with substantial gaps, not a census or a ranking of the world's most profitable monks. S28

The evidence comes in different forms. A seller page establishes what a seller displays. A creator interview establishes what the creator says. An investigation can provide evidence and interpretation while leaving some identities unresolved. A platform counter is a dated observation of a platform counter. None automatically supplies the others.

Throughout the book, **observation** means a finding in the saved research, **reported claim** means a statement attributed to a source, and **interpretation** means our reading of the mechanism. Exercises, budgets, sample scripts, and businesses described as **hypothetical** are invented teaching examples. They are not results from our company or the accounts studied here. The earlier research includes observation notes for two Yang Mun Shorts; the later account expansion checked metadata and pages without completing new moving-image watches. We keep that boundary visible.

Anicca's own avatar factory is a work in progress. At this edition's snapshot, its recovery, end-to-end commercial operation, and sales have not been established. This book does not present a working subscription service, a profitable factory, or a verified hands-off method. It presents public research and a practical design you can test.

Study the structure; leave other people's faces, voices, scripts, and identities with them. Anicca's original book text, scripts, and worksheets are shared under CC BY 4.0 and may be adapted with attribution under that license. Third-party videos, photographs, names, trademarks, and teachings do not become reusable assets because they appear in a bibliography. This book provides no medical treatment or earnings guarantee.

The public companion is the AI Monk Index. Its JSON dataset is a machine-readable interface to the same records. The appendices explain coverage, citations, and reusable worksheets.

Contents

1. The Monk in the Feed
2. What Counts as an AI Monk?
3. A History Without a “First”
4. Yang Mun and Shalev Hani
5. The Creative Grammar of a Monk Short
6. From Scroll to Store
7. The Content Factory
8. Distribution and Measurement
9. Faith, Disclosure, and Health Claims
10. Build a Different Character

Appendices: Working sheets; research method and limits; references.

Chapter 1. The Monk in the Feed

The funniest part of the online monastery is the distance between its atmosphere and its plumbing. Above the fold: a calm face, an open book, candlelight. Below it: a link, a bundle, perhaps a checkout. Stillness has a conversion path.

That observation need not be an accusation. Authors, teachers, artists, and religious communities have material needs. A useful digital book can deserve a price. What matters is the promise attached to it, who makes that promise, and what the audience has been led to believe.

The earlier research notes for Yang Mun's 65-second "Words of Motivation" describe a steady, front-facing scene. The character acknowledges worry, invokes ordinary hopes, and ends with a small reflection for morning and night. There was no visible paid-product pitch in that sample. Its YouTube metadata showed 521,549 views on 21 October 2026. S09

One can see why this is an interesting object. A person anxious about work receives a message addressed to the individual viewer. A familiar visual language suggests patience and distance from ordinary competition. The next swipe would bring another demand, but this voice seems to offer relief from demands. Our interpretation is that the scene makes attention feel like a pause rather than a task. That is an interpretation of a recorded example, not a measured psychological effect.

The next tap can change the relationship. Yang Mun's Journey page displayed a $49.99 program and a $59.99 collection with books at the research snapshot. Now the viewer is a possible customer. S05 The

short video and the offer occupy different moments, yet the recognition created in one can travel into the other. The robe does not need to say “buy” for the profile to have commercial value.

Three questions for one minute

To understand the encounter, ask three questions separately. What experience does the video offer? What identity does the presentation imply? What transaction, if any, follows?

The first question allows room for enjoyment. A fictional character can tell a good story. An animated speaker can prompt a worthwhile reflection. The feeling a viewer reports is not erased by learning how a performance was made. But a moving feeling also does not verify the speaker’s biography or the accuracy of a remedy.

The identity question is narrower. Is this a real teacher? A disclosed fictional character? A real person with a translated or synthetic voice? A compilation account? A temple backdrop provides no answer. The final question concerns the offer: a book, a guide, a supplement, a subscription, a service, or no identified product. The business begins to become legible when those three questions stop borrowing answers from one another.

Consider a hypothetical viewer who watches one 60-second reflection on each of 30 evenings. That is 30 minutes of attention. If the viewer then buys a $15 workbook, the purchase is a separate event. The time does not prove therapeutic benefit, and the purchase does not prove that every previous view caused the sale. This simple separation becomes essential when millions enter the story.

A format with several jobs

The researched profiles do not all lead to the same destination. Yang Mun has a program and bundles. The Zen Storyteller presents a named book in several formats. Yangjun displays a narrowly positioned guide and a separate creator-services listing. Ming San's bio links to an external wellness product. Other records have no verified offer. S05, S19, S22, S23, S16, S28

Calling this a single "AI monk business model" conceals the differences. A book's production and support burden differ from a community's. An external product link may or may not earn commission. A creator-services listing has a business customer rather than the same reader buying a reflection. A useful analysis follows the actual route.

This book therefore studies the monk as a character format: a recognizable presenter, a visual world, a repeated kind of message, and an invitation to do something next. Its cultural power comes partly from what the imagery already means to viewers. Its commercial performance requires evidence farther downstream.

Try a three-column reading

Choose one public post and write three short columns: **experience**, **identity evidence**, **next step**. In the first, describe the feeling or task offered without praising or condemning it. In the second, record an actual disclosure or write "not established." In the third, follow only the public route you can legitimately inspect and name the destination.

For a hypothetical post, the row might read: “A 45-second story about an unfinished chore; disclosed fictional caretaker; free printable planning sheet.” That is more informative than “viral monk account.” It also gives a creator three design problems that can be solved honestly.

Keep a fourth line underneath: **what this record does not show**. It might be product quality, creator identity, sales, or whether the character is synthetic. That line is not administrative clutter. It is where curiosity survives contact with a good-looking page.

Chapter 2. What Counts as an AI Monk?

A robe is a poor detector of software.

An older face, a symmetrical composition, or a strangely smooth voice may invite questions. None, by itself, proves an AI-generated person. A real monk can appear in a stylized video. A fictional monk can be drawn by a human. A synthetic voice can translate a real speaker. A clip can combine all these techniques. The visible result and the production process need separate descriptions.

For this field guide, “AI monk” is a research label for monk-like characters whose synthetic production is claimed or supported by evidence. We also study adjacent accounts because their offers and storytelling illuminate the format. Inclusion does not amount to confirmation.

Five possibilities behind the same screen

A **synthetic persona** is a character presented through generated visual or vocal performance. A **creator-disclosed avatar** may represent a known person or a fictional speaker; the disclosure should say which. A **real religious teacher using a virtual character** has a human teaching identity behind the presentation. **Animation with unresolved production** tells us what the picture looks like without telling us which tools made it. And an **ordinary human creator in religious clothing** may simply be a human creator in religious clothing.

These are analytical possibilities, not labels automatically assigned to everyone in the index. Thai PBS's account of Phra Maha AI is useful precisely because it describes a real Thai monk working with programmers on a 2D character. A digital presentation does not turn the human teacher into a fictional person. S15

The reverse error is equally tempting. If a seller calls a speaker a teacher, that assertion does not settle a dispute about how the character was made. Yang Mun's sources conflict: HeyGen and Columbia describe an AI character associated with Shalev Hani; a first-party statement says the teacher is real and the voice is AI-translated. The index retains `reported`, rather than deciding the case from visual confidence. S01, S02, S03

Status describes evidence, not moral worth

The snapshot uses three status names. `Confirmed` is the highest evidentiary category available in the schema, but no current record meets the archive's application of that category. `Reported` preserves an attributed AI classification with unresolved details or conflicts. `Unverified` means the available record does not establish AI provenance. It does not mean the account has been caught doing anything wrong.

The arithmetic is deliberately unexciting: 4 reported + 71 unverified = 75 records. The 71 include five researched profiles and 66 leads. The nine researched profiles consist of Yang Mun, Ming San, SifuChen California, Leo Chen, Yan Chen/monjechen, Kai Tsukimi, Yangjun, Buddha Would Say, and Master Kaizen's TikTok candidate. Their record quality varies substantially. S28

Yan Chen illustrates the caution. A supplied screenshot showed @monjechen and an AI-generated-profile label. The label was not independently read back in the research. The screenshot's calendar

capture date was unknown. Keeping this record `unverified` preserves both the useful lead and the limit of the evidence. Similar names do not justify merging it with Leo Chen or SifuChen. S29

Ming San illustrates a source conflict rather than a missing label. Maldita includes the profile in its AI health-account investigation; AvatarFactory's commercial directory says public AI provenance is unresolved. The archive preserves that disagreement under `reported`. Its product link and audience count do not decide the identity question. S14, S17

Kai Tsukimi has a named book, a public site, and dated audience metadata. Those observations establish a storytelling brand and a visible offer. They do not establish an AI-generated identity. Buddha Would Say has an illustrated monk format described by a discovery provider, but no first-party production disclosure was established. Master Kaizen has public TikTok metadata, while spoken language and AI provenance remain unresolved. S18, S19, S25, S26

Build a useful record

Begin with the exact account URL, platform, handle, display name, and capture date. Store identity claims in a different field from your classification. Then attach each metric to its platform and source. "1.4 million Instagram followers, rounded connector snapshot, 22 October" is a record. "Huge audience" is an impression.

Keep unknown values empty with a reason. An unavailable follower count is not zero. An English product page does not necessarily establish the language of every video. A profile's country flag is not an operator's residential address. A public-facing character name is not a verified legal identity.

For a hypothetical research budget, suppose a careful first pass takes 12 minutes per candidate. Nine profiles require 108 minutes. Seventy-five require 900 minutes, or 15 hours. The lesson is not that twelve minutes is sufficient; it is that calling all 75 “verified” would conceal a very different amount of work. Record review depth so future researchers can spend their next hour well.

Your minimum useful row is: **account; exact URL; observed date; claimed identity; evidence for AI production; status; offer; unknowns**. Add a conflict field whenever sources disagree. A field guide becomes more valuable when it can say why a label remains provisional.

This discipline leaves room for humor. We can enjoy a synthetic elder explaining the modern inbox without inventing a real elder behind the pixels. The joke works better when the reader knows which part is the character.

Chapter 3. A History Without a "First"

The internet likes an origin story with a founder, a date, and a screenshot. Research often finds a channel date, a later interview, and several missing platforms instead.

The question "Who made the first AI monk?" sounds precise. It hides several possible inventions: the first animated monk, the first real monk using a virtual presenter, the first generated photorealistic character, the first monk account to sell a book, or the first one noticed by a particular language's press. Those are different histories.

This archive cannot resolve a worldwide first. It can build a dated sequence and tell the reader what each date means.

A small chronology, properly labeled

Thai PBS describes Phra Maha AI's collaboration with programmers on a 2D character. Our saved source entry does not establish the article's publication date. It is a relevant virtual-monk precedent, but it should not be forced into a precise chronological slot the record cannot support. S15

A supplied screenshot shows Yang Mun's YouTube channel joining on **4 November 2025**. This dates a channel registration, not the first creation of the persona or its first post anywhere. Public metadata dates the water-themed Short to **10 November 2025**, the health-themed sample to **25 November**, and the motivation sample to **29 November**. These are upload dates for specific items in the saved record. S29, S11, S10, S09

Virtual Humans published its Yang Mun profile on **28 January 2026**. Columbia News Service published its investigation on **24 March 2026**. The first is an industry profile reporting creator claims; the second is independent reporting that includes creator statements. Neither date means the character began on that day. S04, S02

Maldita's **10 July 2026** investigation examined 16 synthetic social profiles. Its **27 July** investigation described a broader set of 250 AI health accounts, with data collected on **20–21 July**. These reports document an ecosystem at specified observation periods. They do not identify a universal starting point for monk avatars. S13, S14

Our **21–22 October 2026** snapshot combines public pages, metadata, prior observation notes, and discovery leads. It expands what can be compared, while still leaving most AI classifications unresolved. A chronology with gaps is more useful than a complete-looking story assembled from guesses.

The date you have is not always the date you want

Imagine a hypothetical character created in February, posted on one platform in March, joined YouTube in May, and received a newspaper profile in August. A search returning only the newspaper would make August look like its birth. A channel screenshot would move that date to May. Neither would uncover March automatically.

That is why a timeline needs at least two date fields: **event date** and **observation date**. Add **date type**: channel join, upload, article publication, page retrieval, or unknown. A screenshot captured in October can display an older upload date; both facts belong in the record.

Counts also have clocks. Yang Mun's channel page displayed 172 videos on 21 October while the Shorts listing returned 143 entries. These are different surfaces, not proof that 29 videos vanished. The difference is 172 − 143 = 29, but arithmetic does not explain the reason for the difference. S08, S12 This is a useful habit: calculate accurately, then stop before the calculation becomes an invented narrative.

A bounded search that can be repeated

To extend a history, define a question narrower than "find everyone." For example: "What is the earliest dated public upload we can establish for this exact account?" Start with its canonical handle, follow its own cross-links, and record the earliest accessible item. Record unavailable surfaces rather than substituting an account with a similar name.

For a hypothetical one-hour session, allocate ten candidates six minutes each. Save the query, language, platform, exact URL, date type, earliest established event, and next missing check. After one hour you have ten auditable attempts, possibly with no new origin date. That result still improves the archive because the next researcher can avoid repeating blind alleys.

Our sample is English-first. Search visibility, discovery recommendations, source language, and accessible tools shape what appears in it. The absence of a documented earlier account in another language says little about whether such an account exists.

History becomes interesting here without a championship belt. The virtual teacher, the generated wellness persona, the illustrated parable, and the direct-to-book brand are converging forms with different obligations. We can trace their intersections without announcing that one robe invented the internet.

Chapter 4. Yang Mun and Shalev Hani

Yang Mun is the central case because the public trail crosses several useful boundaries: a recognizable character, named production technology, a named creator, multiple social surfaces, a direct product offer, large audience claims, and conflicting explanations of identity.

Shalev Hani is the creator named by both HeyGen and Columbia News Service. His public X profile is a creator account, separate from the persona's Instagram @yangmunus and YouTube @Yangmunv2. Columbia reports him as based in Israel. That location belongs to the reported creator biography; it does not locate the character in a real monastery. Exact TikTok handles were not established in this archive. S01, S02, S32, S07, S08

The identity conflict belongs in the case

HeyGen's customer story describes an AI character produced with its tools. Columbia's investigation likewise describes a generated character. A first-party statement linked by Columbia says Yang Mun is a real teacher and only the voice is translated using AI. S01, S02, S03

Those statements cannot be compressed into one clean biography without choosing a conclusion the record does not independently settle. The sites also use different domains: the Journey seller page is on yang-mun.com, while the conflicting identity statement is on yang-muns.com. The similarity of names is not a reason to erase the distinction. Our source register preserves the exact URLs and the relationship reported in the research.

An interview with a production provider is valuable evidence of what that provider and its customer say they did. It also serves the provider's marketing. Independent reporting is a different source type. A first-party response deserves to be represented accurately without being treated as independent proof. Respect for the people involved requires this precision. The unresolved classification is not permission to invent motives or accuse an operator of conduct the evidence does not establish.

An audience with several clocks

The 22 October connector snapshot returned approximately 2.5 million Instagram followers and 433 posts. The YouTube channel page on 21 October displayed about 159,000 subscribers and 172 videos. These figures refer to different platforms and metrics. They should not be added into a count of unique people, because an individual can follow on both platforms. S07, S08

Older successful-looking posts are easier to discover than ordinary ones. The motivation Short in the earlier observation notes showed 521,549 views at capture; the health-themed sample showed 1,102,388; the water-themed metadata entry showed 12,561,496. The last of these was not watched or fact-checked in this research. S09, S10, S11

A later bounded listing adds an important complication. On 22 October the first 20 Shorts entries returned view counts ranging from 12 to 2,300, and 15 titles contained comment-keyword invitations. Publication dates for those entries were not fetched. S12 We cannot calculate equal-age performance, infer a decline, or identify causes. Comparing a recently listed item with an older breakout without its age would be

a poor experiment. The useful finding is narrower: the archive contains both older high-reach examples and a different bounded listing with much smaller snapshot counts.

Two messages, two obligations

The earlier observation record describes the motivation Short as one continuous front-facing shot. Robe, book, candlelight, and temple-like setting hold the visual world together. The message moves from concern about the future to ordinary hopes, then to a repeated reflection. There is no visible paid CTA in that sample. S09

Our interpretation is that this structure reduces the number of things a viewer must understand. There is one speaker, one emotional direction, and one small action. The audience is not asked to learn a complex visual world before following the sentence. This makes an interesting hypothesis for an original production test; the archive does not isolate which element caused the views.

The health-themed Short begins from a sleep concern and moves into a traditional-medicine explanation, according to the earlier notes. That explanation was not independently substantiated. S10 Its greater view count does not make it a better template for claims. A creator can borrow the structural idea of starting with a recognizable problem while choosing a harmless, supportable subject, such as a cluttered desk or a difficult email.

This is where imitation should stop being literal. The transferable question is “How does the script turn a familiar tension into one action?” It is not “Which unverified remedy should I repeat?”

The offer is concrete even when the economics are not

At the saved check, the Journey page displayed a one-time $49.99 price. Its described contents included daily five-to-ten-minute practices, weekly video teaching, a community, and Q&A. A $59.99 Complete Collection added three eBooks, titled *Time To Heal*, *Stillness*, and *Modern Soul*. A separate home-page bundle displayed $19.99. These are seller-listed offers, not independently tested benefits. S05, S06

The $10 difference between $49.99 and $59.99 is a visible offer-design decision. A customer already considering the program sees a relatively small additional price for the books. That is a plausible reason for the structure; the archive does not establish the seller's intention or the bundle's take rate.

Suppose, purely hypothetically, 100 buyers chose between those two prices and 30 selected the higher option. Compared with all 100 buying the $49.99 option, gross receipts would increase by 30 × $10 = $300. This says nothing about actual sales, refunds, costs, or whether the bundle changes how many visitors buy at all. A pricing example becomes useful when its missing variables remain visible.

The seller counters require the same care. The Journey page displayed 30,000-plus readers/books and 2,500 participants; the home page used 20,000-plus readers and a five-million audience statement. Their definitions, periods, platform overlap, and audit trail were unavailable. S05, S06 Multiplying a participant counter by today's listed price would not reconstruct revenue. Participants might have paid different prices, received access another way, or been counted on a different basis. The record does not tell us.

Large numbers that do not reconcile themselves

Columbia reports Hani's statement of more than $300,000 in revenue since launch and a rejected $325,000 acquisition offer. It also reports an AI-character service offered at $799 per hour. S02 Virtual Humans reports a claim of more than $300,000 in profit within 90 days and references a dashboard above $213,000 in gross revenue over three months. S04

A supplied screenshot contains a headline claiming $500,000 in six months. Its exact post URL is not present in the screenshot record. Later search discovery found a creator-post URL, but the body was not read back, so that discovery does not authenticate or reconcile the accounting behind the screenshot. S29, S31

These figures use different periods, categories, and evidence. Revenue is money from sales before relevant deductions. Profit requires a defined treatment of costs. A price for a service is an offer, not evidence of billable hours sold. A proposed acquisition is a possible transaction, not proceeds received. The public record does not supply one reconciled ledger joining these claims.

Nor can we infer deception merely because the figures cannot be reconciled here. We do not know whether the dashboard, interview language, periods, or included business lines differ. The accurate conclusion is an evidence limit: the available material does not let an outside reader compute audited net profit.

To do that work, an analyst would need dated order exports by product, discounts, refunds and chargebacks, payment fees, advertising, production receipts, contractor and labor treatment, support costs, taxes where applicable, and a defined accounting period. To

attribute profit to a particular Short would also require a defensible path from that creative to the order. No robe can substitute for those records.

What a creator can actually take away

HeyGen reports simple scripts, avatar and voice delivery, batching, roughly 20 minutes per video, and increased publishing cadence. The provider describes a sevenfold increase and daily output. Those are case-study assertions, not an independently timed production log. S01

The reusable mechanism is still worth testing: a recognizable original character; a manageable, repeated visual world; one message per asset; a next step related to the message; and a product that continues the same task. The production-time figure is a hypothesis to measure against your own receipts and labor, not a promise to put in a business plan.

Make a one-page case card before copying anything. Record the identity dispute, dated metrics, two observed script structures, visible offer, reported workflow, and missing economics. Underneath, write one original adaptation and one metric it would test. For example: “A disclosed fictional gardener tells a story about unfinished tasks; a free worksheet continues the exercise; measure worksheet visits per post.”

The case becomes more useful as it becomes less magical. It shows a visible route from character to product and a set of creator claims worth examining. It does not turn every quiet voice into a business, or every large counter into a fortune.

Chapter 5. The Creative Grammar of a Monk Short

The shortest route to a dull imitation is to begin with a robe prompt and a proverb generator. You may get something that resembles the category while giving a viewer no reason to remember it.

Begin with the small human situation instead. A person is postponing an email. Someone keeps saving advice without trying it. A reader wants permission to stop improving the evening. These are concrete tensions. A character can enter one of them, make a turn, and leave the viewer with an action that belongs to the story.

Five jobs inside a short script

The examples in the archive suggest five useful script jobs: recognition, tension, turn, action, and destination. This is our practical synthesis, not a universal formula proven across the accounts.

Recognition tells the viewer what situation the video concerns. **Tension** gives the situation movement. **The turn** offers a different way to see or handle it. **The action** is something small enough to attempt. **The destination** says where to go if more help with that exact task is wanted.

Yang Mun's motivation sample supplies an observed movement from worry to a repeatable reflection. The Zen Storyteller supplies caption and offer evidence connecting a life dilemma to a named book. Yangjun's selected caption supplies a direct guide-in-bio invitation.

We have not watched the latter two clips' motion or measured their pacing; they support a message-and-offer comparison, not a shared editing formula. S09, S20, S19, S24, S22

Make the visual world do one job

In the earlier Yang Mun notes, a stable frame gives the speaker room to carry the message. S09 One possible adaptation is a consistent fictional setting with an original object that helps the story: a garden bench, a repair table, an unfinished drawing. Consistency can make production simpler and make recognition easier to test. It need not recreate someone else's temple, face, clothes, or vocal delivery.

Captions should let the viewer follow the actual words. A readable line is more useful than a screen full of animated decoration. A pause can make a turn legible. These are design recommendations for your experiment, not observations of every indexed clip. The point is to choose a few variables you can control, then see whether real viewers understand the piece.

For a hypothetical 60-second script, budget 10 seconds for recognition, 15 for a situation, 15 for a turn, 12 for the action, and eight for a destination. The sum is 60. At an assumed 120 spoken words per minute, uninterrupted speech would allow about 120 words; pauses reduce that allowance. This is a rehearsal budget, not a claim about an account's actual speaking rate. Time the finished performance rather than trusting a word-count estimate.

An original example: the Keeper of Small Tasks

The following fictional character is a caretaker of an imaginary garden, not a real monk or religious authority. Its purpose is gentle comedy about ordinary work. The example is original and may be adapted.

> You have carried that unfinished email through three cups of tea. It has become a very small, very demanding pet.
>
> In my garden, a broken gate does not become easier to repair when I stare at the whole fence. I first find the loose hinge.
>
> Open the draft. Write only the sentence that says what you need. You may dislike the sentence. It still gives you a hinge.
>
> If the email needs more thought, leave a note naming the next question. You have made the work smaller.
>
> There is a blank “next small step” sheet in the profile if you want one.

This script has a familiar problem, a small comic image, one concrete comparison, and an action. Its CTA offers more of the same task. It makes no claim that the exercise treats anxiety or guarantees productivity. A disclosure can appear visibly with the character and in the description: “Fictional AI-assisted storyteller. Original writing by Anicca.” Use that wording only if it accurately describes your production.

Three hooks, one experiment

For that same fictional topic, write three original openings: “The email is still open”; “Your draft has become a pet”; “Before you answer, write the request in one sentence.” The first offers recognition, the second a comic surprise, and the third an instruction. Keep the rest of the script and offer similar enough that you know what changed.

Do not call the winner from one unequal exposure. If a hypothetical first test gives the three versions 200, 2,000, and 150 views, the distribution itself differs. Inspect available retention, clicks, and the kind of response you wanted, and treat the result as a clue. “This opening is worth another try” is often a more defensible conclusion than “This is the viral hook.”

Before producing, read the script aloud. Remove any sentence that introduces a second lesson. Ask whether the action can be understood without the product and whether the product continues the action. Then review the rights and identity of the assets. Originality here is a practical advantage: a character with its own purpose gives you somewhere to go after the first successful post.

The lesson of the archive is not that one script always wins. It is that the path from feeling to action can be studied at the level of actual sentences, while the face and authority that deliver them require separate care.

Chapter 6. From Scroll to Store

Imagine a reader laughing at the email-pet story and tapping the profile. The next page offers an expensive course about becoming rich through avatars. Something has broken. The viewer arrived with one question and met a different business.

A coherent offer continues the work begun in the post. A short parable can lead to a book of original stories. A planning exercise can lead to a workbook. A creator demonstration can lead to a production service. Coherence does not guarantee a purchase, but it gives a visitor a fair chance to understand the transaction.

Follow the destination, not the costume

The researched routes show several different offers:

Case	Publicly documented route	What the record establishes
Yang Mun	Profile to Journey and book collection	Seller-displayed $49.99 and $59.99 offers. S05
Kai Tsukimi / The Zen Storyteller	Recurring storytelling brand to *A Cup of Zen*	Seller-listed Kindle $9.99, paperback $14.99, hardcover $19.99. S19
SifuChen California	Highlighted-story route described by Maldita	Reported €16.45 guide; exact live handle unresolved. S13

Case	Publicly documented route	What the record establishes
Leo Chen	Comment invitation for a book/PDF	Reported lead path; price and paid conversion unknown. S13
Ming San	Bio to an external beetroot product	Visible destination; commission, sales, and efficacy unknown. S16
Yangjun	Guide in bio and a separate services listing	Displayed $29.99 guide and advertised creator services. S22, S23

The table describes routes, not a league table of successful businesses. In particular, The Zen Storyteller's 100,000-plus copies and category-bestseller statement are seller claims; no order ledger or current independently verified Amazon rank was obtained. S19 The account's AI origin also remains unresolved. A strong-looking commercial page can be useful to study without resolving either question.

Yangjun's narrow positioning offers a different lesson. A defined audience and one specific guide make the destination easy to describe. Its medical and symptom-related claims have not been substantiated in this archive, so those claims are not part of the mechanism we recommend adapting. A creator can test the same specificity with an ordinary problem: "A seven-page guide to preparing a difficult meeting" is a clear offer without borrowing medical authority. S22

Choose the smallest product that finishes the task

For the hypothetical Keeper of Small Tasks, start with a printable page rather than a membership. The free page lets the audience try the method. If readers need a sequence of related exercises, an original workbook could become a paid product. Its promise might be: “Ten short exercises for turning an undefined task into a next action.” That sentence describes contents a buyer can inspect. It does not promise a transformed life.

Write the sample before building the checkout. A product with a concrete sample is easier to describe accurately. Decide its file format, length, language, access method, and what help the seller will provide. If a community or ongoing Q&A is included, treat the time required as part of the offer. A digital file can be delivered many times; a human answer still takes someone’s time.

Suppose a hypothetical workbook takes six hours to create, valued at $20 per hour. Its initial labor cost is $120. At a $12 price, assume $1 of transaction-related costs and $1 of support allocation per retained sale, leaving $10 before recovering creation cost. Twelve retained sales would recover that $120 under those assumptions. Advertising, refunds, taxes, and other overhead are excluded, so this is a small contribution calculation rather than a profit forecast. Change the inputs before using it for a real decision.

The calculation helps answer a better question than “How many followers do I need?” It asks how much useful work the product contains, how much delivery costs, and how many actual customers would cover that work. You can then decide whether a pilot is worth its limited budget.

Give each invitation one destination

A direct link is simple to follow, but it still needs a readable page. A comment keyword can reveal interest, but a comment is not consent to every kind of future message. A newsletter can develop a relationship, but it creates an ongoing writing obligation. An affiliate destination may create income, but the relationship and the product's merits require their own evidence.

Pick the route that fits the audience's immediate task. For our fictional caretaker, the first CTA is a free sheet. The sheet can contain a clearly described optional workbook link. A visitor should understand which action downloads the sheet, which action buys a product, and what each will cost. This is an offer-design recommendation, not a statement about any particular platform's current legal requirements.

Before launch, walk the route in ordinary language: "I watched a story, opened the profile, saw a sample, paid for this file, and received it here." If the sentence contains a mysterious leap, fix the page or the product. If the route ends at a displayed checkout, write "checkout available." Save "purchase completed" for an actual completed purchase and "delivered" for a confirmed delivery.

The robe may make a person look. The continuity between a useful post and a well-described product gives that look somewhere sensible to go.

Chapter 7. The Content Factory

The phrase “content factory” suggests a machine that turns a spreadsheet into income while its owner drinks tea. In practice, the factory has rejected renders, a caption that breaks across the wrong words, a missing disclosure, and a product link pointing to yesterday’s idea. The tea is getting cold.

Automation is most useful when the work is already understandable. It can repeat a well-defined operation and preserve records. It cannot decide that an unsupported medical claim is acceptable because the queue is full.

A reported workflow and a proposed one

HeyGen’s Yang Mun story reports simple scripts, avatar and voice delivery, batching, about 20 minutes per video, and increased cadence. GrowwStacks describes an automation route through a content table, an LLM script, HeyGen, captions, and a scheduler. The latter is an automation vendor’s walkthrough, not evidence of Yang Mun’s internal system. Its time-savings claims are not independently measured here. S01, S27

For your own pilot, the proposed route can be shorter: a researched idea, an original script, authorized visual and voice assets, a generated or edited performance, readable captions, a reviewed final file, a public post, and a measured destination. Start with one character, one platform, and one offer. Add machinery only where repetition has become a real problem.

Give each asset a short creative ID. Connect that ID to the script version, source notes, rights basis, final file, platform URL, and product destination. This is ordinary bookkeeping for creative work. It lets you ask “Which version did we actually publish?” without reverse-engineering a folder full of files called final-final.

Make the handoffs visible

A useful production queue has understandable states: **idea**, **script ready**, **render requested**, **render received**, **reviewed**, and **published**. An item cannot become published merely because a scheduler accepted it. The platform’s public URL or authoritative publication receipt completes that handoff.

The same rule matters when a tool times out. A timeout after a render request may mean the request failed, or that it succeeded while the response was lost. Check the provider’s job record before buying a second render. A missing publication response calls for the scheduler or platform’s record before posting again. These are practical ways to avoid duplicates and unnecessary cost, not reasons to stop the work indefinitely.

Keep a named human owner for the final message. The owner checks that the character is described honestly, the script has one point, factual statements have support, assets are authorized, captions match speech, and the destination continues the promise. This review can be brief. It must still be performed on the file that will actually go out.

Count finished assets, including the failures

All numbers in this example are hypothetical. Suppose a month costs $60 in base tools, $24 in generation used in accepted videos, and $36 in failed or rejected generations. Twelve usable videos result. Cash cost is ($60 + $24 + $36) ÷ 12 = **$10 per finished video**. Counting only the $24 would produce a flattering $2 figure that ignores most of the spending.

Now include six hours of editing and review at an assumed $15 per hour. That adds $90, producing an allocated total of $210, or **$17.50 per finished video**. Taxes, promotion, product creation, and other overhead remain outside this example. The purpose is to expose the denominator: a request is not a finished asset, and a finished asset is not necessarily a published one.

For actual production, replace every assumed figure with the relevant invoice, usage receipt, and time record. Define how shared subscriptions are allocated. Record failed attempts once rather than hiding them or charging them to several videos. If unused subscription capacity exists, report both actual cash spending and whatever allocation you use; do not pretend the accounting choice changes the bank balance.

A four-video batch exercise

Prepare four original scripts around one audience problem. Read each aloud before requesting any generation. Reject or revise the weak script while it is still cheap. Generate one sample to test the character, caption size, pronunciation, and disclosure. Then produce the remaining three using what you learned.

Suppose, hypothetically, the batch uses 40 minutes for writing, 20 for revisions, 30 for reviewing renders, and 30 for caption and posting work. That is two hours, averaging 30 minutes per accepted video if all four pass. If two fail, the same two hours produce two assets at an hour each before rework. Record the result you get rather than the one the plan predicted.

Our own operation has not yet earned the right to call this a proven hands-off factory. The design offered here is deliberately testable: inspect the actual file, establish the actual publication, track the actual cost, and preserve uncertain requests until they can be reconciled. Automation should make the next action clearer. A larger queue of unexplained outcomes is only a busier spreadsheet.

Chapter 8. Distribution and Measurement

Your character has been rendered, captioned, and published. A view counter begins to move. This is the moment when a spreadsheet often acquires the powers of a fantasy novel.

Views are useful. They describe a platform-defined exposure measure at a particular time. They do not carry an order receipt inside them. To understand the business, keep following the person and the money through separate, observable events.

Keep the ladder intact

A simple measurement ladder has six stages: **published asset**, **platform response**, **destination visit**, **paid order**, **delivery**, and **retained contribution**. Give each stage its own source.

The public post URL establishes the asset. The platform supplies whatever view, reach, retention, or profile-visit measures are available. Your destination analytics establish visits under their own definitions. The checkout provider supplies completed orders. The delivery system establishes access or dispatch under a defined delivery measure. Refunds, fees, and costs make the last stage possible.

If a stage is unavailable, write “not measured.” A missing visit count does not mean nobody visited. A download request is not always a successful delivery. A paid order later refunded should remain an order in the gross record and a refund in the reconciliation. Deleting inconvenient events gives a prettier sheet and a worse understanding.

For public research, most of the lower ladder is unavailable. The account snapshots in this book provide followers, selected views, public offers, and claims. They do not supply verified per-post orders or net contribution. S28 A private operator could measure more in its own business, but should not publish personal customer details to make the graph look convincing.

A small funnel, fully labeled

Here is a hypothetical pilot. One original post records 20,000 views during a defined observation window. Its dedicated link records 160 product-page visits. Eight completed $12 orders follow that tagged route. One is refunded. Assume $4 in total transaction costs, $3 in delivery costs, and $40 in allocated production cost for the post. These are invented inputs, not current fees from any provider.

Measure	Calculation	Hypothetical result
Visits per view	160 ÷ 20,000	0.8%
Paid orders per visit	8 ÷ 160	5%
Gross receipts	8 × $12	$96
Receipts after refund	$96 − $12	$84
Contribution after listed costs	$84 − $4 − $3 − $40	$37

The $37 excludes labor beyond the production allocation, advertising, taxes, and other overhead. It is not net profit. The attribution window is a chosen rule, not proof that the post alone caused each purchase. Readers may encounter several posts, switch devices, share a link, or arrive later without the tag.

That does not make the measurement worthless. It makes its claim specific: “Eight completed orders used this recorded route during this window.” You can act on that finding while preserving uncertainty about causation.

Give the creative a durable name

Use a creative ID such as `garden-email-01`. Record platform, public URL, publication time, script version, hook variant, CTA, offer version, and link tag. Add dated metric snapshots rather than overwriting the only number. The name should refer to the actual asset, because “Tuesday’s video” becomes ambiguous when you publish on several platforms.

Track cash costs and labor separately. Store receipts securely in your own operation and publish only the appropriate totals. Reconcile orders against the payment provider’s authoritative status. A dashboard that says “checkout started” is a different source from one that says “paid.” If the order succeeded but delivery did not, repair delivery under that order rather than charging the buyer again.

Test a question you can answer

“Can I become viral?” is a wish, not an experiment. “Does a comic opening send more qualified readers to the same free sheet than a direct instruction?” is a testable question.

In a small pilot, keep the character, audience topic, destination, and product promise stable while changing one meaningful script element. Choose an observation window before comparing. Note the exposure

each asset actually received. A post with more clicks may simply have had more views; visits per view can answer a different question from total visits.

Do not manufacture certainty from tiny numbers. If one hypothetical version gets two orders from 40 visits and another gets one from 30, both results are sparse. A sensible next action may be to try the clearer message again with another original story. There is no need to crown a permanent winner.

The October research offers a warning about windows. Yang Mun's older high-view clips and the first 20 later listing entries do not have comparable ages in the available record. Master Kaizen's indexed and fetched view figures also differ within the saved snapshot. Preserve source-specific counts rather than averaging them into an invented exact value. S09, S11, S12, S26

At the end of each pilot, write four sentences: what was published; what was observed; what remains unmeasured; what one change follows. If no orders are established, write that. If there are orders but no delivery confirmation, name that boundary. The most useful report points to a next action without pretending it has already happened.

Chapter 9. Faith, Disclosure, and Health Claims

An invented gardener can complain about an inbox. An invented religious authority can appear to speak for a tradition, a lineage, and a community that never appointed it. The difference begins before the script's first word.

Robes, temple settings, names, and accents carry meaning. They are not a neutral collection of visual prompts. A creator borrowing them should ask what a reasonable viewer is likely to infer: actual ordination, personal experience, historical knowledge, medical expertise, or simply a clearly fictional story.

Make the speaker understandable

A useful disclosure answers ordinary questions. Is the character fictional? Which parts are synthetic or AI-assisted? Who prepares the words? Is a real teacher represented, and with whose permission? Does the seller have a commercial relationship with a product?

For a wholly fictional production, a plain description might be: "An original fictional character, performed with AI-assisted visuals and a licensed synthetic voice. Scripts are written and reviewed by the publisher." This is a hypothetical template. Change it to fit the truth; do not claim human review or a license you do not have.

Put the information where viewers can reasonably encounter it, including the video or its immediate description and the profile. Platform labels can supplement that explanation. This book does not document every platform's current policy or every jurisdiction's require-

ments. Check the current official rules for the actual service and location before publishing; do not turn a sentence here into a claim of legal compliance.

The identity dispute around Yang Mun demonstrates why the words matter. “Real teacher with AI translation” and “AI-generated character” tell the viewer different things about the speaker. The point of preserving the conflict is to make interpretation possible, not to turn an unresolved identity into entertainment at a person’s expense. S01, S02, S03

Treat traditions as sources, not props

If you attribute a teaching, identify its source and distinguish a translation, a paraphrase, and your own reflection. If you do not know whether a saying comes from a particular teacher or text, do not place that name under it. A freshly generated sentence does not acquire an ancient pedigree because it sounds calm.

The virtual presentation of a real monk described by Thai PBS shows another possibility: a human teaching identity using a digital form. S15 That case should not be collapsed into the fictional-influencer category. The teacher’s role, the production method, and the representation of the tradition each deserve their own description.

Maldita’s wider health-account investigation discusses an “ancestral wisdom” group containing monks and other stereotyped figures. Its 82 accounts in that group are not 82 Buddhist monks. S14 The finding is a reason to examine how imagery substitutes for expertise, not permission to repeat the same stereotype with a better renderer.

An original character can have texture without claiming a culture's authority. Give it specific habits, a purpose, and limits. Our fictional garden caretaker knows where it left the rake and admits when a story is invented. It does not claim access to secret medicine, represent all Buddhists, or pretend its biography happened.

Health changes the stakes

A viewer may enjoy a harmless story and also be vulnerable to an authoritative-sounding remedy. A health claim requires its own qualified evidence; popularity cannot provide that evidence. The researched examples include unsupported or unverified claims, and this book studies their presentation rather than endorsing them. S10, S14, S16, S22

For a creator without the necessary expertise and substantiation, the practical choice is to stay with a different topic. A story about preparing a question for an appointment differs from advising a viewer how to diagnose or treat a condition. An evidence-free "ancient remedy" is not improved by adding a soothing soundtrack or a disclaimer beneath it.

Review the actual claim

Use a small claim ledger. Copy each factual assertion from your own draft, name its source, state what that source supports, and decide whether to keep, narrow, or remove it. Treat health, income, and identity statements as substantive claims even if they appear only in a caption or sales-page heading.

For a hypothetical batch of 12 scripts, assume five minutes of basic review each: 60 minutes. If four contain assertions requiring another 15 minutes of source work each, allow another hour. The two-hour estimate is a planning assumption. If a claim cannot be substantiated within the available work, remove it or choose a different subject. Do not keep it because generation has already been paid for.

Respect can improve the creative premise. A character who openly inhabits fiction can be funny, surprising, and useful without inviting viewers to mistake costume for credentials. Transparency gives the audience the information needed to choose its relationship with the story.

Chapter 10. Build a Different Character

You now have enough information to make an imitation. A more useful next step is to make a character whose reason for existing is your own.

The fictional Keeper of Small Tasks began with a narrow audience situation: people carrying an unfinished ordinary task around all day. Its comic personality follows from that job. It mistakes drafts for pets, loses a rake, and prefers a loose hinge to a philosophical explanation of the entire fence. Its visual world is an imaginary garden. Its first useful object is a sheet that helps someone name a next action.

No part requires a false monastery, an unlicensed human likeness, or a miraculous promise. The character can still be odd. Oddness is allowed to be an artistic choice rather than an identity claim.

Days 1–2: choose a person and a task

Write one sentence: “This character helps [specific audience] with [ordinary task] through [kind of story].” Then write the limit: “It does not claim [authority or outcome].” For our example, the audience is people facing small unfinished work; the story is a gentle comic comparison; the limit excludes religious and medical authority.

Study three relevant public examples before deciding how to present the idea. The archive’s comparison of Yang Mun, Kai Tsukimi, and Yangjun separates a recorded one-point reflection, a story-to-book

route, and a narrow guide offer. These examples have different evidence quality and unresolved operators; they are not three independently proven profitable factories. S09, S19, S22

Write down what you will adapt: a clear opening situation, one point, a stable original setting, or continuity from story to product. Write down what stays with the original creator: words, face, voice, biographical identity, and protected media. This is a design brief you can actually use.

Days 3–4: make the useful object

Create the free sheet or a substantive product sample first. Fill it out yourself for a real ordinary task you are willing to discuss. If you cannot explain how it helps, a character will not fix it.

The Keeper's sheet asks four questions: What is unfinished? What is the next visible action? What information is missing? When will I revisit it? Its promise is small and inspectable. A paid workbook, if one is later produced, must add concrete material rather than repackage a free sentence into a large claim.

Describe any offer in plain terms: contents, format, language, price, delivery, support, and limitations. At this stage an unpublished product remains unpublished. Do not write testimonials, sold counters, or earnings headlines to make the page feel occupied.

Days 5–7: produce a small batch

Write six original scripts with two opening styles across three situations. Use the recognition–tension–turn–action–destination structure as a drafting aid, not a rule every sentence must obey. Read them aloud, remove unsupported claims, and check asset rights and disclosures.

Produce one sample before the other five. Check the actual output for comprehensibility: can a viewer read the captions, follow the story, and understand who the character is? Revise the production settings that interfere with those tasks. Then finish the batch and record each accepted asset's cost and time.

For a hypothetical pilot, set a **$90 cash cap**: up to $60 for production, $15 for delivery-related needs, and $15 reserve. These are allowances, not quoted service prices or money already spent. Use no paid advertising in this example. Separately value four hours of labor at an assumed $20 per hour, or $80. Full use of the cash allowance plus that labor would total $170; actual expenses may be lower. Stop adding paid generations when the cap is reached and simplify the work.

Days 8–12: publish and observe

Use one platform initially. Publish the reviewed assets on a cadence you can actually maintain, keeping the public URL, creative ID, disclosure, and destination together. Confirm publication through the platform or scheduler's authoritative record. Inspect the route to the free sheet or real product, because a beautiful post pointing to an unavailable destination cannot complete its job.

Choose the main measure before looking for a winner. For a free-sheet pilot, it might be completed sheet downloads per recorded visit. For a real paid product, it might be retained orders and their contribution after specified costs. Keep view and engagement figures as useful context. Avoid changing the product, price, character, topic, and hook simultaneously if you want to understand the next result.

A hypothetical six-post pilot with no established sales is not automatically a failure or evidence of zero demand. First inspect the measured boundary: were the posts public, were the links usable, did visits occur, did the offer communicate its contents, and was payment available? Where data is absent, improve the relevant observation. Where the message is weak, change the message. Where the product is missing, make the product.

Days 13–14: decide from the record

Reconcile the pilot's costs, publication records, visits, orders if any, refunds, and delivery outcomes. Write a short report using Appendix A. Name the strongest observed result and the most important missing fact. Choose one next change within the budget you can justify.

You may discover that people enjoy the stories but want a different worksheet. You may find a useful product with an uninteresting presenter. You may find that production takes too long. Those are different results and suggest different work. More posts are not the answer to every question.

Keep the work portable

Preserve a simple production recipe, original scripts, asset permissions, source notes, and a table connecting creative IDs to outcomes. These records make later automation easier. The public JSON companion can support research tools without becoming a second, contradictory database: its `accounts`, `sources`, status fields, and dated snapshots should travel together. Its existence is not a claim that a downloadable, proven factory or subscription application is available. S28

If you eventually publish reusable code, remove secrets and private customer records, document setup and limitations, and license only material you control. Share the workflow that actually worked, including its failed steps and costs. Let an open-source promise follow a reproducible result rather than substitute for one.

The art in this subject is not limited to making a convincing face. It is the art of choosing a human situation, giving a character an honest role in it, and following the work all the way to the reader. A small truthful story has room to be funny. It also has somewhere to go after the scroll.

Appendix A. Working sheets

Copy these sheets for your own research or production. Blank fields mean “not yet recorded,” never zero. Do not put credentials or private customer details in a public sheet.

A1. Evidence card

Field	Your entry
Exact account / platform / URL	_____
Observation date and acquisition method	_____
Public-facing identity claim	_____
AI-production evidence and source URL	_____
Status and reason	_____
Conflicting statement and exact source	_____
Metric / value / platform / snapshot date	_____
Visible offer / listed price / currency	_____
What the evidence establishes	_____
What remains unknown / next check	_____

A2. Original character and script brief

Audience: _____ Ordinary situation: _____

Character's role: _____ Limits of its authority: _____

Disclosure: _____ Asset rights basis: _____

Opening recognition: _____ Tension: _____

One turn: _____ One action: _____

Destination and what it actually contains: _____

Reference mechanism being adapted, with URL: _____

Factual claims and support, or claims removed: _____

Read-aloud duration / final-file review / reviewer: _____

A3. Publication and effect ledger

Field	Your entry
Creative ID / script and offer versions	_____
Requested publication time / actual time	_____
Platform publication record / public URL	_____
Request uncertain? What readback is needed?	_____
Destination / link tag / observation window	_____
Dated views / visits / other measured events	_____
Completed orders / refunds / delivery status	_____

Field	Your entry
Evidence location for each event	_____
Owner / next action	_____

A4. Economics and pilot review

Cash paid: base tools ____ + accepted usage ____ + failed usage ____ + other ____ = ____

Finished assets: ____ Published assets: ____ Labor hours and valuation: ____

Orders × actual price: ____ Refunds: ____ Fees: ____

Delivery, support, promotion, and allocated production: ____

Contribution under this definition: ____ Excluded costs: ____

Observed result: _____ Missing measurement: _____

One next change: _____ Budget and owner: _____

Appendix B. Research method and coverage limits

This first edition uses Anicca's saved research from 21–22 October 2026 and the public sources registered with it. The public companion contains 75 account records and 48 source entries. Nine profiles received research beyond initial discovery; 66 remain discovery leads. AI classification is reported for four and unverified for 71, with none confirmed under the current archive's application of its categories. S28

The initial five-profile study combined reporting, seller pages, platform metadata, supplied screenshots, and earlier observation notes for two Yang Mun Shorts. The later English-first expansion added four profiles and enriched existing records using a connected public-research provider, public site retrieval, and yt-dlp metadata. It completed no new moving-image watches. A caption, thumbnail description, or discovery summary is not relabeled as a watched video.

Supplied screenshots are leads with limitations. Some show a channel date, a label, a follower count, or a cropped product link. Their calendar capture dates are not established unless explicitly recorded. The screenshots themselves are not reproduced here. S29 points to the public dataset's descriptive source entries, so readers can inspect what was transcribed while recognizing that the underlying images are not publicly available.

Discovery began partly from supplied screenshots and followed public references. It is neither random sampling nor exhaustive search. English is prioritized, other languages are undercovered, some live handles remain unresolved, and independence between operators is

not established. Similar names are not merged without evidence. Country flags, seller-declared locations, and platform location signals do not establish an operator's residence.

Metrics retain platform, source, and date. Rounded connector counts remain rounded. Cross-platform followers are not added into unique people. Missing fields remain unknown. The first 20 Shorts listing entries are a bounded feed sample without fetched publication dates, not a complete performance series. Older high-view examples are selected cases, not estimates of typical output.

Maldita's 250-profile study concerns a broader AI health-account sample. Its 53.7-million-plus follower total, 45,600-plus posts, seven-language coverage, and 82-account ancestral-wisdom group are the investigation's reported findings. They are not monk-only market totals and should not be added to our sample. S14

First-party products and counters establish displayed offers and seller assertions. No private order ledger, refund export, verified royalties, per-creative sales, or audited profit was obtained for the studied accounts. The book makes no market-size, causal growth, or guaranteed-income estimate. Its hypothetical calculations are teaching devices; its original scripts do not reproduce creator scripts.

Exploratory direct Instagram retrieval stopped after review of its access notice; this book does not offer that route as a continuing authorized collector. S30 The appropriate next research depends on current access permissions and usable official or connected sources. New findings should preserve acquisition method and uncertainty rather than overwrite history.

Anicca's original book text, scripts, worksheets, and the companion dataset's Anicca-authored compilation and summaries use CC BY 4.0. Third-party content retains its own rights. The source manifest ac-

companying this manuscript preserves the 48-entry register and maps this book's citations to it. Dates in the references are saved observation dates, not a claim that the linked pages remain unchanged when you read them.

Appendix C. References

Reference labels link to the original public source, except S29, which links to the public descriptions of non-public supplied screenshots. Named bylines were not consistently preserved in the saved register; publisher or account attribution is used rather than inventing an author. Creator assertions are attributed as such. The companion source manifest records the exact mapping and evidence limits.

Identity, creator, and direct offers

S01. HeyGen, Yang Mun customer story. Provider/customer account of Shalev Hani's production process; creator and performance claims, not an independent audit. English page checked 22 October 2026; Japanese version checked 21 October. Source

S02. Columbia News Service, "Millions found comfort in a Buddhist monk, but he was never real." Published 24 March 2026; checked 21 October. Independent reporting containing Hani's identity, revenue, acquisition-offer, and service-price statements. Source

S03. Yangmuns.com, identity statement concerning *Time To Heal*. Checked 21 October 2026; publication date unestablished. First-party real-teacher/AI-translated-voice assertion that conflicts with S01 and S02. Source

S04. Virtual Humans, Yang Mun profile. Published 28 January 2026; checked 21 October. Industry profile reporting creator audience and financial claims; no audited ledger supplied. Source

S05. Yang-mun.com, Journey seller page. Checked 21 October 2026. $49.99 program, $59.99 collection, contents and counters are seller displays. Source

S06. Yang-mun.com, home page. Checked 21 October 2026. Separate $19.99 bundle and counters without a reconciled definition. Source

S32. Shalev Hani, public X profile. Checked 21 October 2026. Creator-profile reference; existence of a profile does not verify financial claims. Source

Platform records and creative samples

S07. Yang Mun, Instagram @yangmunus. Connector snapshot, 22 October 2026: rounded 2.5 million followers and 433 posts. Source

S08. Yang Mun, YouTube @Yangmunv2. Public channel check, 21 October 2026: approximately 159,000 subscribers and 172 videos. Source

S09. Yang Mun, motivation Short. Uploaded 29 November 2025; metadata and earlier observation notes, checked 21 October 2026: 65 seconds, 521,549 views. Source

S10. Yang Mun, wisdom/health Short. Uploaded 25 November 2025; metadata and earlier observation notes, checked 21 October 2026: 78 seconds, 1,102,388 views. Health content unsubstantiated here. Source

S11. Yang Mun, water-themed Short. Uploaded 10 November 2025; metadata only, checked 21 October 2026: 96 seconds, 12,561,496 views. No completed watch or health fact-check. Source

S12. Yang Mun, public Shorts listing. Checks on 21–22 October 2026. Different listing surfaces and bounded samples; 22 October first-20 sample lacks fetched per-entry publication dates. Source

S16. Ming San, Instagram @mingsanhealing. Connector snapshot, 22 October 2026: rounded 1.4 million followers, 175 posts, external product link. Source

S17. AvatarFactory, Ming San profile. Checked 21 October 2026. Commercial directory that leaves AI origin unresolved, conflicting with Maldita's classification. Source

S18. Kai Tsukimi, Instagram @thezenstoryteller. Connector snapshot, 22 October 2026: rounded 369,900 followers and 971 posts. Source

S19. The Zen Storyteller, *A Cup of Zen* seller page. Checked 22 October 2026. Prices, 21-story description, copies-sold counter, and rank claim are seller statements. Source

S20. Kai Tsukimi, selected pinned Reel. Uploaded 25 July 2026; metadata checked 22 October. Indexed rounded 2.8 million plays; no moving-image watch. Source

S21. Yangjun, Instagram @monkyangjun. Connector snapshot, 22 October 2026: rounded 95,600 followers and 323 posts. Source

S22. Yangjun, guide seller page. Checked 22 October 2026. $29.99 eight-week menopause guide and destination link; health benefits unverified. Source

S23. Yangjun, Collabstr seller listing. Checked 22 October 2026. Self-described wellness-character services and location; no verified partnerships or orders. Source

S24. Yangjun, guide-in-bio Reel. Uploaded 21 October 2026; metadata checked 22 October. Indexed rounded 2,100 plays; caption evidence only, no moving-image watch. Source

S25. Buddha Would Say, selected Reel. Uploaded 18 September 2026; metadata/index checked 22 October. Rounded 1.1 million views and 9,400 followers; AI origin unverified. Source

S26. Master Kaizen, selected TikTok. Uploaded 6 October 2026; metadata/index checked 22 October. Source-specific view counts differ; language and AI origin unresolved. Source

Investigations, method, and adjacent workflow

S13. Maldita.es, 16-profile investigation. Published 10 July 2026; checked 21 October and revisited 22 October. Reports synthetic profiles and Leo Chen/SifuChen offer routes, not verified sales. Source

S14. Maldita.es, 250 AI health-profile investigation. Published 27 July 2026; data collected 20–21 July; checked 21 October. Broader health sample with stated methodology, not a monk census. Source

S15. Thai PBS World, Phra Maha AI profile. Checked 21 October 2026; publication date not established in the saved record. Real monk using a 2D virtual character. Source

S27. GrowwStacks, avatar-video automation walkthrough. Published 14 January 2026; checked 21 October. Vendor-proposed workflow and time-savings claims; no independently verified Yang Mun system. Source

S28. Anicca, AI Monk Index and open dataset. Snapshot updated 22 October 2026. Original compilation; nine researched profiles, 66 leads, 48 sources; provisional classifications and linked evidence. Index · JSON dataset

S29. Supplied screenshots, descriptive register. Reviewed 21 October 2026. Underlying images are not public. Source IDs cover Yang Mun's channel date, Ming San, monjechen, Hani's headline, and following-list leads; capture dates and some destinations are unknown. Public descriptions

S30. Instagram, robots/access notice. Checked 22 October 2026. Describes restrictions encountered by the research route at that snapshot. Source

S31. Shalev Hani X post, discovery URL. Found 22 October 2026 through search. Post body not read back; retained as a discovery lead, not verified earnings evidence. Discovery URL